

benchmarking_nvidia_rtx_2080_200603

June 3, 2020

1 Benchmarking numpy / scikit-image / scipy vs clesperanto

```
[1]: import clesperanto as cle
import numpy as np
import time
import matplotlib.pyplot as plt

num_iterations = 10

# measure execution time of a given method
def benchmark(function, kwargs):
    times = []
    for i in range(0, num_iterations):
        start_time = time.time()
        function(**kwargs)
        delta_time = time.time() - start_time
        times = times + [delta_time]
        # print(delta_time)

    # return median of measurements to ignore warmup-effects
    return np.median(times)

def benchmark_size(method_np, method_cle, method_cle_alloc):
    times_ref = []
    times_cle = []
    times_cle_alloc = []
    sizes = []
    for size in [1, 2, 4, 8, 16, 32, 64]:

        input1 = np.zeros((1024, 1024, size))
        cl_input1 = cle.push(input1)
        cl_input2 = cle.create(cl_input1.shape)

        time_ref = benchmark(method_np, {"image":input1})
```

```

        time_cle = benchmark(method_cle, {"image":cl_input1, "output":
↪cl_input2})
        time_cle_alloc = benchmark(method_cle_alloc, {"image":cl_input1})

        times_ref = times_ref + [time_ref]
        times_cle = times_cle + [time_cle]
        times_cle_alloc = times_cle_alloc + [time_cle_alloc]
        sizes = sizes + [size]

    plt.plot(sizes, times_ref, 'r--', sizes, times_cle, 'g--', sizes,
↪times_cle_alloc, 'b--');
    plt.ylabel('Time / ms')
    plt.xlabel('Image size / MB')
    plt.legend(("ref", "cle", "cle+alloc"));
    plt.show()

    print("\nSizes (MB)          " + str(sizes))
    print("Times ref (s)          " + str(np.round(times_ref, 4)))
    print("Times cle (s)          " + str(np.round(times_cle, 4)))
    print("Times cle+alloc (s)    " + str(np.round(times_cle_alloc, 4)))

```

1.1 Thresholding

```


[2]: # RED: thresholding of a numpy array
def threshold_ref(image):
    thresholded = image > 100
    return thresholded

# GREEN: thresholding of a pre-existing opencl array (no push, pull or alloc)
def threshold_cle(image, output):
    cle.greater_constant(image, output, 100)

# BLUE: allocate result memory + thresholding
def threshold_cle_alloc(image):
    thresholded = cle.create(image.shape)
    cle.greater_constant(image, thresholded, 100)

benchmark_size(threshold_ref, threshold_cle, threshold_cle_alloc)

```



Sizes (MB)	[1, 2, 4, 8, 16, 32, 64]
Times ref (s)	[0.0005 0.001 0.003 0.006 0.0125 0.023 0.039]
Times cle (s)	[0.0035 0.003 0.0035 0.003 0.004 0.003 0.003]
Times cle+alloc (s)	[0.004 0.004 0.005 0.006 0.007 0.009 0.0145]

1.2 Gaussian blur radius 2

```
[3]: from skimage.filters import gaussian


radius = 2

def gaussian_blur_filter_ref(image):
    filtered = gaussian(image, sigma=radius)
    return filtered

def gaussian_blur_filter_cle(image, output):
    cle.gaussian_blur(image, output, radius, radius, radius)

def gaussian_blur_filter_cle_alloc(image):
    filtered = cle.create(image.shape)
    cle.gaussian_blur(image, filtered, radius, radius, radius)


benchmark_size(gaussian_blur_filter_ref, gaussian_blur_filter_cle, ↵
↵gaussian_blur_filter_cle_alloc)
```



Sizes (MB)	[1, 2, 4, 8, 16, 32, 64]
Times ref (s)	[0.051 0.1581 0.534 0.7797 0.8848 1.1981 3.818]
Times cle (s)	[0.008 0.009 0.01 0.013 0.02 0.0335 0.052]
Times cle+alloc (s)	[0.008 0.009 0.011 0.015 0.022 0.034 0.0581]

1.3 Gaussian blur radius 10

```
[4]: radius = 10
benchmark_size(gaussian_blur_filter_ref, gaussian_blur_filter_cle,
↳ gaussian_blur_filter_cle_alloc)
```



Sizes (MB)	[1, 2, 4, 8, 16, 32, 64]
Times ref (s)	[0.1231 0.2858 0.7614 1.8947 2.2588 3.1704 6.4054]
Times cle (s)	[0.009 0.011 0.015 0.023 0.038 0.0621 0.1146]
Times cle+alloc (s)	[0.009 0.0115 0.016 0.023 0.0371 0.0661 0.1176]

1.4 Binary erosion


```
[5]: from skimage.morphology import binary_erosion

def binary_erosion_ref(image):
    filtered = binary_erosion(image)
    return filtered

def binary_erosion_cle(image, output):
    cle.erode_box(image, output)

def binary_erosion_cle_alloc(image):
    filtered = cle.create(image.shape)
    cle.erode_box(image, filtered)

benchmark_size(binary_erosion_ref, binary_erosion_cle, binary_erosion_cle_alloc)
```



Sizes (MB)	[1, 2, 4, 8, 16, 32, 64]
Times ref (s)	[0.006 0.011 0.0205 0.039 0.0841 0.1591 0.3163]
Times cle (s)	[0.0035 0.003 0.003 0.003 0.0035 0.003 0.003]
Times cle+alloc (s)	[0.004 0.004 0.004 0.005 0.007 0.01 0.0145]

1.5 Mean filter radius=2

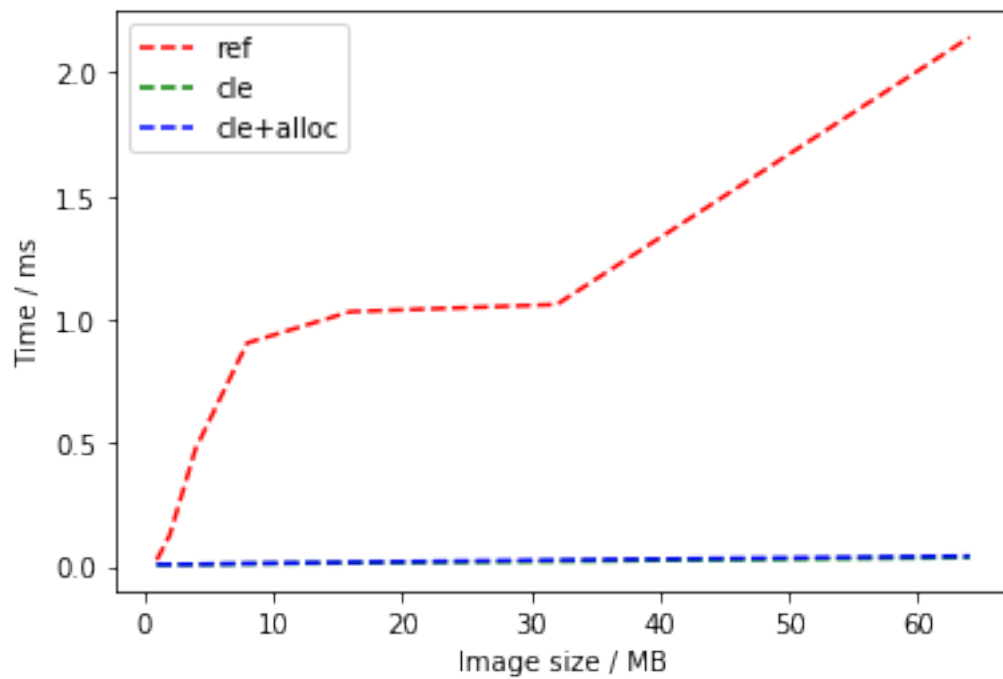
```
[6]: import scipy.ndimage.filters as spf

radius = 2
def mean_filter_ref(image):
    # todo: not sure if size is a radius or a diameter. Check documentation
    # https://docs.scipy.org/doc/scipy/reference/generated/scipy.ndimage.
    ↪uniform_filter.html#scipy.ndimage.uniform_filter
    filtered = spf.uniform_filter(image, size=radius)
    return filtered

def mean_filter_cle(image, output):
    cle.mean_box(image, output, radius, radius, radius)

def mean_filter_cle_alloc(image):
    filtered = cle.create(image.shape)
    cle.mean_box(image, filtered, radius, radius, radius)
```

```
benchmark_size(mean_filter_ref, mean_filter_cle, mean_filter_cle_alloc)
```



```
Sizes (MB)      [1, 2, 4, 8, 16, 32, 64]
Times ref (s)    [0.027  0.1261 0.4752 0.9048 1.0319 1.061  2.142 ]
Times cle (s)    [0.008  0.009  0.009  0.012  0.015  0.022  0.037]
Times cle+alloc (s) [0.009  0.009  0.01  0.013  0.018  0.027  0.043]
```

[]: